



UPPSALA
UNIVERSITET

Department of Information Technology

Scientific Computing for Data Analysis

Davoud Mirzaei

Lecture 2: Monte Carlo Method - II

Agenda

- ▶ Well-known distributions
- ▶ Monte Carlo integration
- ▶ Inverse transform method (ITM)

General structure of Monte Carlo (MC) algorithm

```
Input N : (Number of observations)  
for  $k = 1 : N$   
    perform one stochastic simulation/process  
     $result[k] = \text{result of the simulation}$   
end  
 $FinalResult = \text{mean}(result) \text{ or other statistical calculation}$ 
```

- ▶ MC works with random numbers/stochastic (random) processes
- ▶ Random numbers are generated (sampled) from different distributions



Some well-known probability distributions

Uniform distribution: We write $X \sim \mathcal{U}(a, b)$, the pdf of X is

$$f(x) = \frac{1}{b-a} \begin{cases} 1, & x \in [a, b] \\ 0, & x \notin [a, b] \end{cases}$$

The cdf of X is

$$F(x) = \mathbb{P}(X \leq x) = \int_{-\infty}^x f(s)ds = \begin{cases} 0, & x < a \\ \int_a^x f(s)ds = \frac{x-a}{b-a}, & x \in [a, b] \\ 1, & x > b \end{cases}$$




What is the area under the pdf graph?

Some well-known probability distributions


Normal distribution: We write $X \sim \mathcal{N}(\mu, \sigma^2)$, and read X is a normal variable with mean μ and variance σ^2 (standard deviation σ). The pdf of X is

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right), \quad x \in (-\infty, \infty) \quad (1)$$

Plots of pdf $f(x)$ and cdf $F(x)$ for $\mu = 0$ and two different values for σ :



pdf of normal dist.



cdf of normal dist.

Some well-known probability distributions



Exponential distribution: We write $X \sim \mathcal{Exp}(\lambda)$, its pdf f is given by

$$f(x) = \lambda e^{-\lambda x}, \quad x \in [0, \infty), \quad \lambda > 0$$

and its cdf F by

$$F(x) = \int_0^x \lambda e^{-\lambda s} ds = 1 - e^{-\lambda x}, \quad x \geq 0.$$

The mean or expectation of X is $\mu = \frac{1}{\lambda}$ and its variance is $\sigma^2 = \frac{1}{\lambda^2}$.




Models things like waiting times, time between earthquakes, time between calls

Some well-known probability distributions

Discrete distributions:

- ▶ Previous distributions are all examples of continuous distributions
- ▶ There are also some discrete distributions, Example: dice,

$$f(x) = \begin{cases} \frac{1}{6}, & x \in \{1, 2, 3, 4, 5, 6\} \\ 0, & \text{otherwise} \end{cases}$$




What does the cdf look like?

Some well-known probability distributions


General **Discrete Distributions**: $X \sim \mathcal{DD}([x_1, \dots, x_m], [p_1, \dots, p_m])$

$$f(x_k) = p_k, \quad k = 1, 2, \dots, \quad \sum_{k=1}^{\infty} p_k = 1.$$

x_k	x_1	x_2	\dots	x_m
p_k	p_1	p_2	\dots	p_m



$$F(x) = \sum_{k: x_k \leq x} p_k$$



Random numbers in NumPy

Use `rand` and `randn` to sample from uniform and normal distributions, respectively:

```
# a uniform random number in interval [0,1)
x = numpy.random.rand()

# a (N x 1) array of uniform random numbers in [0,1)
x = numpy.random.rand(N,1)

# a uniform random number in interval [a,b)
x = (b-a)*numpy.random.rand() + a

# a standard normal random number (mu = 0, sigma = 1)
x = numpy.random.randn()

# a (N x 1) array of standard normal numbers (mu = 0, sigma = 1)
x = numpy.random.randn(N,1)

# a normal random number with mean mu and standard deviation s)
x = mu + s*numpy.random.randn()
```



Note: If $X \sim \mathcal{U}(0, 1)$ then $a + (b - a)X \sim \mathcal{U}(a, b)$
If $X \sim \mathcal{N}(0, 1)$ then $\mu + \sigma X \sim \mathcal{N}(\mu, \sigma^2)$

Plot histograms (pdf and cdf)

Example: Normal samples

```
import numpy as np
import matplotlib.pyplot as plt
N = 5000
mu, s = 5, np.sqrt(3)
X = mu + s*np.random.randn(N,1)
# histogram plot (pdf)
plt.figure()
plt.hist(X, bins=50, histtype='bar', edgecolor='black', density='true')
plt.title('pdf, normal samples  $\mu=5, \sigma^2 = 3$ ')
plt.xlabel('$x$'); plt.ylabel('Frequency $\%$')

# histogram plot (cdf)
plt.figure()
plt.hist(X, bins=50, histtype='bar', edgecolor='black', density='true', cumulative='true')
plt.title('cdf, normal samples  $\mu=5, \sigma^2 = 3$ ')
plt.xlabel('$x$'); plt.ylabel('Frequency $\%$')
```



Expectation and Variance of a random variable

Definition: Assume that X is a random variable with pdf f . The *expectation* or *mean* of X is defined by

$$\text{Discrete : } \mu = \mathbb{E}[X] = \sum_k x_k p_k = \sum_k x_k \mathbb{P}(X = x_k)$$

$$\text{Continuous : } \mu = \mathbb{E}[X] = \int x f(x) dx$$

The variance of X is defined as

$$\sigma^2 = \text{Var}(X) = \mathbb{E}[(X - \mu)^2]$$

“Expecting how much X is deviated from the mean!”

Connection to integration: If $X \sim f(x)$ is a continuous random variable and g is a function, then $g(X)$ is another random variable. The expectation of $g(X)$ is obtained as

$$\mathbb{E}[g(X)] = \int g(x) f(x) dx$$

Exercise

Ex 1: Compute the mean and variance of a 6-sided dice drawing:

x_k	1	2	3	4	5	6
p_k	1/6	1/6	1/6	1/6	1/6	1/6

$$\mu = \mathbb{E}[X] = \sum_{k=1}^6 x_k p_k = \dots, \quad \sigma^2 = \sum_{k=1}^6 (x_k - \mu)^2 p_k = \dots$$

Ex 2: Compute the mean and variance of a random variable $X \sim \mathcal{U}(a, b)$

$$\mu = \mathbb{E}[X] = \int_a^b x f(x) dx = \dots, \quad \sigma^2 = \mathbb{E}[(X - \mu)^2] = \int_a^b (x - \mu)^2 f(x) dx = \dots$$

Ex 3: Compute the mean and variance of a random variable $X \sim \mathcal{Exp}(\lambda)$

$$\mu = \mathbb{E}[X] = \int_0^{\infty} x f(x) dx = \dots, \quad \sigma^2 = \int_0^{\infty} (x - \mu)^2 f(x) dx = \dots$$

Monte Carlo integration

Assume that we aim to estimate the generic integral

$$\int_a^b g(x)f(x) dx$$

where f is a density function for random variable X . The function g is called the *performance* function.

Key point: The above integral is just $\mathbb{E}[g(X)]$, the mean of variable $g(X)$, thus **Monte Carlo method** can be applied to estimate it:

- ▶ Generate samples x_1, \dots, x_N from pdf f
- ▶ Compute observations $g(x_k)$ and set

$$\bar{g}_N = \frac{1}{N} \sum_{k=1}^N g(x_k)$$

as an estimate for the integral.

- ▶ Random points are generated from f , only g appears in front of the summation symbol

Example: Consider the 1D integral

$$\int_0^1 g(x) dx$$

This integral is indeed the expectation of $g(X)$ for **uniform** variable X on $[0, 1]$:

$$\mathbb{E}[g(X)] = \int_0^1 g(x) \cdot 1 dx, \quad X \sim \mathcal{U}(0, 1)$$

- **MC method:** Choose N uniformly distributed random points x_1, x_2, \dots, x_N in $[0, 1]$



Then

$$\int_0^1 g(x) dx \approx \frac{1}{N} [g(x_1) + g(x_2) + \dots + g(x_N)]$$

Approximate the value of integral $\int_0^1 \sin x \, dx$ using MC method with different number of points $N = 10^k$, $k = 0, 1, 2, 3, 4, 5$

```
def g_fun(x):  
    return np.sin(x)  
  
int_exact = 1-np.cos(1)      # exact value of integral for comparison  
int_mc = np.zeros([6,1])  
for k in range(6):  
    N = 10**k  
    X = np.random.rand(N,1)  # generate N uniform random points in [0,1]  
    g = g_fun(X)  
    int_mc[k] = np.sum(g)/N   # take mean  
print('int_mc = \n',np.abs(int_exact-int_mc))
```

The output for an execution is:


```
int_mc =  
[[0.149659 ]  
 [0.07784648]  
 [0.0130975 ]  
 [0.00131985]  
 [0.00172056]  
 [0.0005902 ]]
```

A new execution leads to a different result, but still shows a slow convergence

Monte Carlo integration (uniform distr.)

Estimate the integral of g on finite interval $[a, b]$

$$\int_a^b g(x)dx = (b-a) \int_a^b g(x) \frac{1}{b-a} dx \approx \underbrace{(b-a)}_{\text{width}} \underbrace{\frac{1}{N} \sum_{k=1}^N g(x_k)}_{\text{mean height}}, \quad x_k \in \mathcal{U}(a, b)$$



Monte Carlo integration (uniform distr.)

An example in Python: Estimate the value of integral

$$\int_{-\pi/2}^{\pi/2} x^2 \cos x \, dx$$

using MC method and $N = 1000$

```
import numpy as np
def g_fun(x):          # g(x) defined in function g_fun
    return x**2*np.cos(x)

a,b = -np.pi/2,np.pi/2  # integral bounds a and b
N = 1000                  # number of realizations
result = np.empty(N)
for k in range(N):
    u = np.random.rand() # uniform point u in [0,1]
    x = (b-a)*u + a       # uniform point x in [a,b]
    result[k] = g_fun(x)  # function evaluation
int_mc = (b-a)*np.mean(result) # Monte Carlo estimation
```

Note: If $U \sim \mathcal{U}(0, 1)$ then $X = (b - a)U + a \sim \mathcal{U}(a, b)$

Monte Carlo integration (example with normal distribution)

To estimate the value of integral

$$I = \int_{-\infty}^{\infty} (x^4 - x + 1)e^{-x^2/2} dx$$

using MC method, we can write

$$I = \sqrt{2\pi} \int_{-\infty}^{\infty} \underbrace{(x^4 - x + 1)}_{g(x)} \underbrace{\frac{1}{\sqrt{2\pi}} e^{-x^2/2}}_{f(x)} dx$$

$f(x)$ is the pdf of $\mathcal{N}(0, 1)$ on $(-\infty, \infty)$, so

$$I \approx \sqrt{2\pi} \times \frac{1}{N} \sum_{k=1}^N (x_k^4 - x_k + 1)$$


where x_k are generated from $\mathcal{N}(0, 1)$.

Compared to a deterministic method


As a deterministic method, consider the mid-point (MP) rule for integration:

$$\begin{aligned}\int_0^1 g(x)dx &= h [g(x_1^*) + g(x_2^*) + \cdots + g(x_N^*)] + \mathcal{O}(h^2) \\ &= \frac{1}{N} [g(x_1^*) + g(x_2^*) + \cdots + g(x_N^*)] + \mathcal{O}(N^{-2})\end{aligned}$$

where $h = 1/N$ and N is the number of integration points in the interval $[0, 1]$, and $x_k^* = (x_{k-1} + x_k)/2$ are mid points (**equidistance points, not random**).



Stochastic vs. Deterministic



Comparing the error and **order of convergence** in a log-log plot:

- ▶ Mid-point method (deterministic): $p = 2$, $e = \mathcal{O}(N^{-2}) = c \cdot N^{-2}$
- ▶ Monte Carlo method (stochastic): (ten simulations, different results), $p = 0.5$, $e = \mathcal{O}(N^{-0.5}) = c \cdot N^{-0.5}$


Stochastic vs. Deterministic

- ▶ The order of convergence of mid-point method is 2. Error behaves as $c \cdot h^2 = c \cdot N^{-2}$ in the **maximum norm**.
- ▶ The MC method converges in **probability** with convergence rate 0.5 **independent of the dimension**, the error behaves as $c \cdot N^{-1/2}$
- ▶ Which one is better? It depends on the dimension:
 - ▶ In 1D $h = \mathcal{O}(N^{-1})$ so mid-point error is $\mathcal{O}(N^{-2})$
 - ▶ In 2D $h = \mathcal{O}(N^{-1/2})$ so mid-point error is $\mathcal{O}(N^{-1})$
 - ▶ In 3D $h = \mathcal{O}(N^{-1/3})$ so mid-point error is $\mathcal{O}(N^{-2/3})$
 - ▶ In 4D $h = \mathcal{O}(N^{-1/4})$ so mid-point error is $\mathcal{O}(N^{-1/2})$
 - ▶ In 5D $h = \mathcal{O}(N^{-1/5})$ so mid-point error is $\mathcal{O}(N^{-2/5})$
- ▶ For dimensions greater than 4 the MC has a faster convergence

Answer the question when the mid-point rule is replaced by the Simpson rule (convergence rate h^4)?

Conclusion: Monte Carlo is a better choice for integration in higher dimensions!

Comparing h and N in 1D and 2D





- ▶ In 1D $N = 10$ and $h = N^{-1} = 0.1$
- ▶ In 2D $N = 100$ and $h = N^{-2} = 0.1$

In deterministic methods for integration, to keep the same accuracy in 2D the number of points must be squared. In 3D it must be cubed, ...

Approximate π using a 2D MC method

Since the area of a circle of radius 1 is π so we can generate N uniformly distributed random points in square $[-1, 1]^2$ and count the points inside the circle:

$$\frac{\text{\#points inside circle}}{\text{\#points inside square}} \approx \frac{\text{area of circle}}{\text{area of square}} = \frac{\pi}{4}$$



How is it connected to MC integration?

Generating random points from other distributions

Inverse transform method (ITM)


Assume that we want to generate X from a distribution $f(x)$. Let the cdf be denoted by $F(x)$. If F is invertible and $U \sim \mathcal{U}(0, 1)$ then

$$X = F^{-1}(U) \sim f$$

Why? Since F is invertible and $\mathbb{P}(U \leq u) = u$, we have

$$\mathbb{P}(X \leq x) = \mathbb{P}(F^{-1}(U) \leq x) = \mathbb{P}(U \leq F(x)) = F(x)$$

This means that the cdf of X is F or the pdf of X is f .




Animation: Watch

The definition of inverse function: If F is continuous and increasing then F^{-1} has the usual definition

$$F^{-1}(y) = \{x : F(x) = y\}$$

To cover all cases including discrete and nondecreasing functions we have the following definition:

$$F^{-1}(y) = \min\{x : F(x) \geq y\}, \quad 0 \leq y \leq 1.$$



Generate random points from the pdf below using ITM:

$$f(x) = \begin{cases} 2x, & x \in [0, 1] \\ 0, & \text{otherwise,} \end{cases}$$

Solution: first we compute the corresponding cdf. It is enough to consider f on its support, i.e. for $x \in [0, 1]$

$$F(x) = \int_0^x f(s)ds = \int_0^x 2s ds = s^2 \Big|_0^x = x^2$$

Then we calculate the inverse of F which is $F^{-1}(x) = \sqrt{x}$. (write $y = F(x)$, switch x and y , $F(y) = x$ i.e. $y^2 = x$, or $y = \sqrt{x}$) Then we generate a uniform variable U and finally we set

$$X = F^{-1}(U) = \sqrt{U}$$

```
U = np.random.rand(N,1)
X = np.sqrt(U)
```

Sampling from exponential distribution using ITM

If $X \sim \mathcal{Exp}(\lambda)$, then $f(x) = \lambda e^{-\lambda x}$ for $x \in [0, \infty)$ and its cdf F is computed as

$$F(x) = \int_0^x \lambda e^{-\lambda s} ds = 1 - e^{-\lambda x}, \quad x \geq 0.$$

The inverse of F is (why?)

$$F^{-1}(x) = -\frac{1}{\lambda} \ln(1 - x)$$

To sample from the exponential distribution, we assume $U \sim \mathcal{U}(0, 1)$ and set

$$X = -\frac{1}{\lambda} \ln(1 - U) \sim \mathcal{Exp}(\lambda)$$

```
def RandExp(lam,N):  
    # lam: distribution parameter, N: number of requested samples  
    U = np.random.rand(N) # generate N uniform numbers in [0,1)  
    X = -1/lam*np.log(1-U) # use inverse transform to generate X  
    return X
```

Sampling from Bernoulli distribution

We say X has Bernoulli distribution with probability $p \in [0, 1]$ and write $Ber(p)$ if

$$\mathbb{P}(X = 0) = p, \quad \mathbb{P}(X = 1) = 1 - p$$

So Bernoulli is the simplest discrete distribution with $x = \{0, 1\}$ and probability vector $\{1 - p, p\}$.

x	0	1
$f(x)$	$1 - p$	p

Example: coin flip, $p = 0.5$ if coin is fair

Sampling: Generate a uniform variable $U \sim \mathcal{U}(0, 1)$. If $U \leq p$ set $X = 1$ otherwise $X = 0$.

```
def RandBer(p,N):  
    X = np.zeros(N)  
    U = np.random.rand(N)  
    idx = np.where(U <= p)  
    X[idx] = 1  
    return X
```


Sampling from a general discrete distributions

Assume that X has a discrete distribution $\mathcal{DD}([x_1, \dots, x_m], [p_1, \dots, p_m])$

x_k	x_1	x_2	\dots	x_m
p_k	p_1	p_2	\dots	p_m

Let $x_1 < x_2 < \dots$. The cdf of X is simply given by

$$F(x) = \sum_{k: x_k \leq x} p_k.$$



Sampling from discrete distributions

ITM: according to definition of F^{-1} , first generate a uniform distribution $U \sim \mathcal{U}(0, 1)$ and then find the smallest positive integer k such that $U \leq F(x_k)$. Finally $X = x_k$ is reported. Equivalently

$$X = \begin{cases} x_1, & 0 \leq U < p_1 \\ x_2, & p_1 \leq U < p_1 + p_2 \\ x_3, & p_1 + p_2 \leq U < p_1 + p_2 + p_3 \\ \vdots & \vdots \end{cases}$$

```
def RandDisct(x,p,N):  
    # x: sorted states  
    # p: probabilities  
    # N: number of requested samples  
    cdf = np.cumsum(p)          # compute the cumulative vector  
    U = np.random.rand(N)       # generate N uniform numbers in [0,1)  
    idx = np.searchsorted(cdf, U) # search U values in cdf intervals  
    return x[idx]               # return corresponding states
```

Example: A three-dice rolling game

Consider a game where you toss 3 six-sided dice simultaneously.



If there are no doubles or triples then you win the sum of the three dice. Otherwise you lose all you have won so far. Let the **random variable X be the value of your winning after 10 tosses**. Here comes an example:

toss	outcome	winnings	toss	outcome	winnings
1	3, 2, 1	6	6	5, 4, 6	25
2	4, 1, 2	13	7	2, 2, 2	0
3	3, 5, 3	0	8	3, 5, 1	9
4	1, 6, 1	0	9	1, 4, 4	0
5	2, 5, 3	10	10	1, 5, 2	8 $=: x_1$

What is the expected gain? Solve with Python.

Monte Carlo solution to the 3-dice game

```
# Monte Carlo simulation for the 3-dice game
x = np.array([1,2,3,4,5,6])
p = np.array([1/6,1/6,1/6,1/6,1/6,1/6])
N = 10**4
X = np.empty(N)
for k in range(N):
    win = 0
    for j in range(10):
        Dice = RandDisct(x,p,3)      # 3-dice rolling simulation
        if len(np.unique(Dice)) < 3: # check for double or triple
            win = 0
        else:
            win += np.sum(Dice)
    X[k] = win
print('Expected Win =', np.mean(X))
```

A single execution gives:

```
Expected Win = 13.0534
```


An application of MC to estimate probabilities

Assume that X is a random variable and we want to estimate the probabilities like $\mathbb{P}(X \leq a)$ or $\mathbb{P}(X = a)$,

Example: In the 3-dice game estimate the probability that your winning will be greater than \$20. It means estimation of $\mathbb{P}(X \geq 20)$.

We just need to count the the number of outcomes greater than 20 and divide it by total outcomes:

```
pr_atleast20 = np.size(np.where(X >= 20))/N
```

A run of above code for $N = 10^4$ results in $\mathbb{P}(X \geq 20) = 0.2603$. So the probability of winning at least \$20 is near 0.26%

For grade 3:

Algorithm2_20230819

The task is to approximate the value of integral

$$\int_0^{\infty} (1+x) \exp(-2x) dx$$

using the Monte Carlo method on five random points

0.0108, 0.0602, 0.3568, 0.8921, 1.7759

which are exponentially distributed according to probability density function (pdf) $f(x) = 2 \exp(-2x)$. What is the approximate value?

Select one alternative:

☐ 0.80958

☐ 1.61916

☐ 0.59823

☐ 0.29911

For higher grades:

Grade45_2_20230819

Assume that you are running a lumber mill in Krokomb and you are trying to **estimate the production price** of a single piece of your standard framing timber. The cost to produce your standard framing timber includes labor, energy and trees. Assume the cost of labor is constant at 3 SEK per piece of framing timber. The cost of energy needed to make a single piece of framing timber is normally distributed with mean $\mu_E = 0.5$ SEK and standard deviation $\sigma_E = 0.1$ SEK. The price of tree needed to make a single piece of timber is distributed according to Weibull pdf $f_T(x) = 5x^4 \exp(-x^5)$ for $x \in [0, \infty)$.

Design a Monte Carlo algorithm to estimate the mean production price p_{mean} and variance $p_{variance}$ of standard framing timber.

Assume the function `randn()` exists and that it returns one standard normal number (with mean 0 and variance 1) every time that it is called. However, it is necessary to provide a detailed formulation of how we can generate a random number from the given Weibull distribution and incorporate it into the Monte Carlo algorithm.

Fill in your answer here or write in the answer sheet and hand it in.

 [Help](#)

Format



B

I

U

\times_2

\times^2

I_x

