



UPPSALA
UNIVERSITET

Department of Information Technology

Scientific Computing for Data Analysis

Davoud Mirzaei

Block 2

Numerical Linear Algebra- I

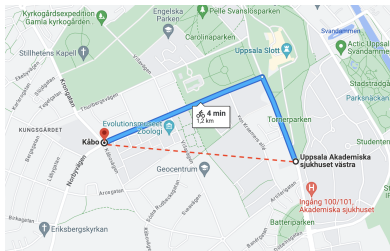
Lecture 5: Review of norms + regression analysis and least squares

Agenda

- ▶ Vector and matrix norm
- ▶ Stability and condition numbers
- ▶ Polynomial fitting
- ▶ Least squares problem and normal equation

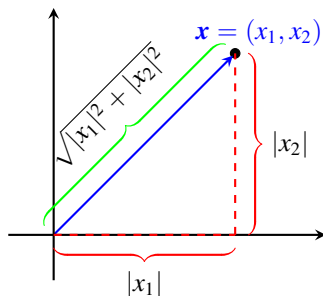
Norms and condition numbers

- What is a norm? Norm means distance:



- Norms for vectors and matrices - Why?
Problem: Measure size of vectors or matrices. What is “small” and what is “large”?
Problem: Measure distance between vectors or matrices. When are they “close together” or “far apart”?
Answers are given by norms
- Also: Tool to analyze convergence and stability of algorithms

Vector norms



Some frequently used vector norms: Assume that $\mathbf{x} = (x_1, \dots, x_n)$ is an n -vector:

- ▶ $\|\mathbf{x}\|_2 = \sqrt{|x_1|^2 + \dots + |x_n|^2}$ (norm 2 or Euclidean norm)
- ▶ $\|\mathbf{x}\|_1 = |x_1| + |x_2| + \dots + |x_n|$ (norm 1 or sum norm)
- ▶ $\|\mathbf{x}\|_\infty = \max_{1 \leq k \leq n} \{|x_k|\}$ (norm infinity or maximum norm)

In 1D all the norms are identical with $|x|$ (the absolute value of x)

Assume that \mathbf{x}, \mathbf{y} are two vectors and α is a real number. Any norm must possess the following properties:

- ▶ $\|\mathbf{x}\| \geq 0$ and $\|\mathbf{x}\| = 0 \iff \mathbf{x} = \mathbf{0}$ (positivity)
- ▶ $\|\alpha \mathbf{x}\| = |\alpha| \|\mathbf{x}\|$ (homogeneity)
- ▶ $\|\mathbf{x} + \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\|$ (triangle inequality)

We can simply show that norm 1 and norm infinity satisfy these conditions. For norm 2 the only challenging part to prove might be the triangle inequality.

Note: Norm 2 is induced from the inner product

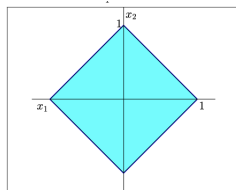
$$\langle \mathbf{x}, \mathbf{y} \rangle := \mathbf{x}^T \mathbf{y} = x_1 y_1 + x_2 y_2 + \cdots + x_n y_n$$

It is clear that

$$\|\mathbf{x}\|_2 = \sqrt{\langle \mathbf{x}, \mathbf{x} \rangle}$$

Vector norms

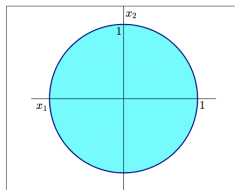
The unit ball $B = \{x : \|x\| \leq 1\}$ in different norms:



$$\|x\|_1 \leq 1$$

or

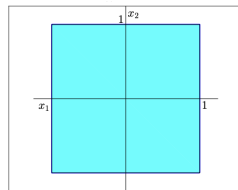
$$|x_1| + |x_2| \leq 1$$



$$\|x\|_2 \leq 1$$

or

$$(|x_1|^2 + |x_2|^2)^{1/2} \leq 1$$



$$\|x\|_\infty \leq 1$$

or

$$\max\{|x_1|, |x_2|\} \leq 1$$

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

For matrices we again look for some norm definitions with the same properties:

- ▶ $\|A\| \geq 0$ and $\|A\| = 0 \iff A = \mathbf{0}$ (positivity)
- ▶ $\|\alpha A\| = |\alpha| \|A\|$ (homogeneity)
- ▶ $\|A + B\| \leq \|A\| + \|B\|$ (triangle inequality)

and usually with two additional properties:

- ▶ $\|A\mathbf{x}\|_* \leq \|A\| \|\mathbf{x}\|_*$ for some vector norm $\|\cdot\|_*$ (consistency)
- ▶ $\|AB\| \leq \|A\| \|B\|$ (submultiplicativity)

These additional properties are sometimes necessary because when dealing with products, matrices differ from mn -vectors!

Matrix norms

Example: Let $\|A\|_{\max} := \max_{i,j} |a_{ij}|$. It satisfies the first three properties but (at least) not the last property. For example

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \quad \text{then} \quad AB = \begin{bmatrix} 3 & 3 \\ 7 & 7 \end{bmatrix}$$

We see that $\|A\|_{\max} = 4$, $\|B\|_{\max} = 1$, $\|AB\|_{\max} = 7$ thus

$$\|AB\| > \|A\| \|B\|$$

Example: But if we define $\|A\|_F := \sqrt{\sum_{i,j} |a_{ij}|^2}$ then it satisfies all 5 properties. (The fourth property with $\|x\|_* = \|x\|_2$). This norm is called the **Frobenius norm** (It is not norm 2 of a matrix).

$$A = \begin{bmatrix} 5 & -4 & 2 \\ -1 & 2 & 3 \\ -2 & 1 & 0 \end{bmatrix}, \quad \|A\|_F = \sqrt{25 + 16 + 4 + 1 + 4 + 9 + 4 + 1 + 0} = 8$$

Example (continued): Norm Frobenius is consistent with the Euclidean vector norm (norm 2): $\|A\mathbf{x}\|_2 \leq \|A\|_F \|\mathbf{x}\|_2$ for all vectors \mathbf{x} of size n .

$$A = \begin{bmatrix} 5 & -4 & 2 \\ -1 & 2 & 3 \\ -2 & 1 & 0 \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} -3 \\ 1 \\ -1 \end{bmatrix}$$

$$A\mathbf{x} = \begin{bmatrix} 3 \\ 12 \\ 0 \end{bmatrix} \quad A\mathbf{y} = \begin{bmatrix} -21 \\ 2 \\ 7 \end{bmatrix}$$

$$\|A\mathbf{x}\|_2 = \sqrt{153}, \quad \|A\mathbf{y}\|_2 = \sqrt{494}, \quad \|\mathbf{x}\|_2 = \sqrt{14}, \quad \|\mathbf{y}\|_2 = \sqrt{11}, \quad \|A\|_F = 8$$

$$\blacktriangleright \sqrt{153} = \|A\mathbf{x}\|_2 \leq \|A\|_F \|\mathbf{x}\|_2 = 8\sqrt{14}$$

$$\blacktriangleright \sqrt{494} = \|A\mathbf{y}\|_2 \leq \|A\|_F \|\mathbf{y}\|_2 = 8\sqrt{11}$$

Inequality holds true not only for these two vectors but also for all vectors of size 3×1 . (we skip the math. proof!)

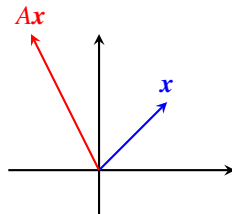
Matrix norms

Natural norms

We can define a class of matrix norms via vector norms. If we apply a matrix A on a vector x we get a new vector Ax perhaps with a different length and a different direction:

$$A = \begin{bmatrix} -2 & 1 \\ 1 & 1 \end{bmatrix}, \quad x = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

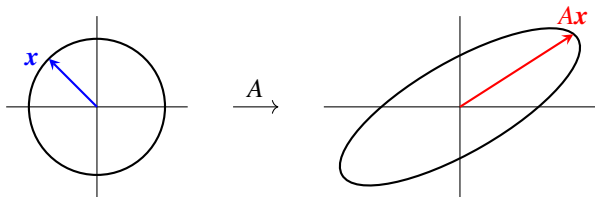
$$Ax = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$$



- Matrix A scales (stretches or shrinks) and rotates x

Matrix norms

Consider a set of all vectors x with $\|x\| = 1$ (e.g. in norm 2: all vector with heads on the unit circle)



The norm of A is identified by measuring “how much A stretches the unit vectors?”

The definition:

$$\|A\|_p := \max_{x: \|x\|=1} \|Ax\|_p, \quad p = 1, 2, \dots, \infty$$

We can prove that such definition is indeed a matrix norm which satisfies all 5 properties above.

This norm is called p -norm or natural norm of matrix A

Matrix norms

It is impossible to compute the p -norm of a matrix by directly using the definition above. However, from the definition we can prove that

- ▶ $\|A\|_1 = \max_{1 \leq j \leq n} \sum_{i=1}^m |a_{ij}|$
- ▶ $\|A\|_\infty = \max_{1 \leq i \leq m} \sum_{j=1}^n |a_{ij}| = \|A^T\|_1$
- ▶ $\|A\|_2 = \sqrt{\lambda_{\max}(A^T A)}$, (λ_{\max} maximum eigenvalue)

Example:

$$A = \begin{bmatrix} 0 & -1 \\ 5 & -3 \\ -2 & 1 \end{bmatrix}, \quad \begin{aligned} \|A\|_1 &= \max\{7, 5\} = 7 \\ \|A\|_\infty &= \max\{1, 8, 3\} = 8 \\ \|A\|_2 &=? \end{aligned}$$

$$\begin{aligned} A^T A &= \begin{bmatrix} 29 & -17 \\ -17 & 11 \end{bmatrix} \Rightarrow \lambda_1 \doteq 0.7646, \lambda_2 \doteq 39.2354 \\ \Rightarrow \|A\|_2 &\doteq \sqrt{39.2354} \doteq 6.2638 \end{aligned}$$

The matrix p norm is consistent with the vector p norm (why?):

$$\|A\mathbf{x}\|_p \leq \|A\|_p \|\mathbf{x}\|_p, \quad p = 1, 2, \dots, \infty$$

Also we can show (how?)

$$\|AB\|_p \leq \|A\|_p \|B\|_p$$

Some Exercises:

- ▶ Show that $\|\mathbf{x}\|_\infty \leq \|\mathbf{x}\|_2$ for all vectors \mathbf{x} .
- ▶ Show that $\|\mathbf{x}\|_\infty \leq \|\mathbf{x}\|_1$ for all vectors \mathbf{x} .
- ▶ If A is symmetric then $\|A\|_1 = \|A\|_\infty$
- ▶ If A is symmetric then $\|A\|_2 = \max\{|\lambda_k(A)|\}$
- ▶ Show that $\|A\|_F = \sqrt{\text{trace}(A^T A)}$ (The trace of a matrix is the sum of its diagonals)

The code for norms of vectors and matrices are provided in `numpy.linalg.norm` library:

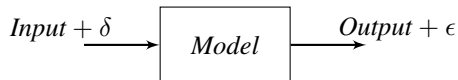
```
import numpy as np

x = np.array([5,-2,1])           # x of size 3x1
A = np.array([[1,2,3],[4,-1,-2],[0,3,2]]) # A of size 3x3

Norm1_x = np.linalg.norm(x,1)    # norm 1 of x
Norm2_x = np.linalg.norm(x,2)    # norm 2 of x
NormInf_x = np.linalg.norm(x,np.inf) # norm infinity of x

Norm1_A = np.linalg.norm(A,1)    # norm 1 of A
Norm2_A = np.linalg.norm(A,2)    # norm 2 of A
NormInf_A = np.linalg.norm(A,np.inf) # norm infinity of A
NormFro_A = np.linalg.norm(A,'fro') # norm Frobenius of A

Norm_x = np.linalg.norm(x) # (default) norm 2 of x
Norm_A = np.linalg.norm(A) # (default) norm Frobenius of x
```



- ▶ The input data of a model (problem) is always subjected to some perturbations (like measurement errors, roundoff errors)
- ▶ The fundamental question then is:
How much the solution of the perturbed problem is close to the solution of the original problem?
- ▶ If the sensitivity is low, the model is referred to as **well-conditioned**; otherwise, it is termed **ill-conditioned**
- ▶ If we replace the term “model” with “method” then then we say the method is either **stable** or **unstable**

Conditioning of linear systems

Example: Consider $Ax = b$ for $n \times n$ nonsingular matrix A and n -vector b . Let A is Hilbert matrix:

$$A = \begin{bmatrix} 1 & \frac{1}{2} & \cdots & \frac{1}{n} \\ \frac{1}{2} & \frac{1}{3} & \cdots & \frac{1}{n+1} \\ \vdots & \vdots & & \vdots \\ \frac{1}{n} & \frac{1}{n+1} & \cdots & \frac{1}{2n-1} \end{bmatrix} \quad (1)$$

and $x = [1, 1, \dots, 1]^T$ and $b = Ax$. Now given that b find x in Python:

```
import numpy as np
import scipy as sp
n = 11
A = sp.linalg.hilbert(n)
b, x = np.sum(A, axis = 0), np.ones(n)
x_hat = np.linalg.solve(A,b)
print("x_hat = ", x_hat)
err = np.linalg.norm(x-x_hat,np.inf)/np.linalg.norm(x,np.inf)
print('RelativeError = ', err)
```

Example: conditioning of linear systems

We expect a solution \hat{x} close to exact solution $x = [1, 1, \dots, 1]^T$ but the output is:

```
x_hat = [0.99999999 1.00000071 0.99998152 1.0002066  
         0.99876902 1.00432858 0.99057385 1.01285236  
         0.98932278 1.00494066 0.99902392]
```

```
RelativeError = 0.01285235836299159
```

- ▶ Compare the relative error to the machine epsilon $\epsilon_{ps} \approx 10^{-16}$, fifteen significant digits are lost!
- ▶ Here $n = 11$ is small. Much more terrible results for larger n
- ▶ What is the source of this disaster? The problem $Ax = b$ itself or the algorithm of `solve` function in Python?
- ▶ The algorithm is stable but the system $Ax = b$ with A Hilbert matrix is ill-conditioned, small perturbation in inputs (here A and b) results in a huge error in the solution (here x). The error was amplified by a factor of 10^{15}

Condition number of a matrix

- ▶ Consider the system $Ax = b$ and the perturbed system $A(x + \varepsilon) = b + \delta$ we get $\varepsilon = A^{-1}\delta$. (Assuming perturbation only in b for simplicity)
- ▶ Take norm from both sides:

$$\|\varepsilon\| = \|A^{-1}\delta\| \leq \|A^{-1}\| \|\delta\|$$

- ▶ From original system $Ax = b$ we have $\|b\| \leq \|A\| \|x\|$, or

$$\frac{1}{\|x\|} \leq \frac{\|A\|}{\|b\|}$$

- ▶ Multiplying both sides of recent equations we get

$$\frac{\|\varepsilon\|}{\|x\|} \leq \|A\| \cdot \|A^{-1}\| \frac{\|\delta\|}{\|b\|}$$

- ▶ At the worst case the error in the output is amplified by $\|A\| \cdot \|A^{-1}\|$. So the **condition number** of A is defined as

$$\text{cond}(A) := \|A\| \cdot \|A^{-1}\|$$

Example:

$$A = \begin{bmatrix} 2 & 4 & -1 \\ 2 & 5 & 2 \\ -1 & -1 & 1 \end{bmatrix} \xrightarrow{\text{Verify!}} A^{-1} = \frac{1}{5} \begin{bmatrix} -7 & 3 & -13 \\ 4 & -1 & 6 \\ -3 & 2 & -2 \end{bmatrix}$$

Thus

$$\text{cond}_1(A) = \|A\|_1 \cdot \|A^{-1}\|_1 = 10 \cdot \frac{21}{5} = 42$$

$$\text{cond}_\infty(A) = \|A\|_\infty \cdot \|A^{-1}\|_\infty = 9 \cdot \frac{23}{5} = 41.4$$

$$\text{cond}_2(A) = \|A\|_2 \cdot \|A^{-1}\|_2 \doteq 7.1561 \times 3.4184 \doteq 24.4621$$

Condition number of a matrix

- ▶ In the Hilbert example the roundoff error in \mathbf{b} is of order machine epsilon (eps), i.e.

$$\frac{\|\delta\|_{\infty}}{\|\mathbf{b}\|_{\infty}} = \text{eps} \approx 10^{-16}$$

The condition number of a Hilbert matrix of size 11×11 is approximately 10^{+15} . So we have

$$\frac{\|\mathbf{x} - \hat{\mathbf{x}}\|_{\infty}}{\|\mathbf{x}\|_{\infty}} \lesssim 10^{+15} \cdot 10^{-16} = 0.1$$

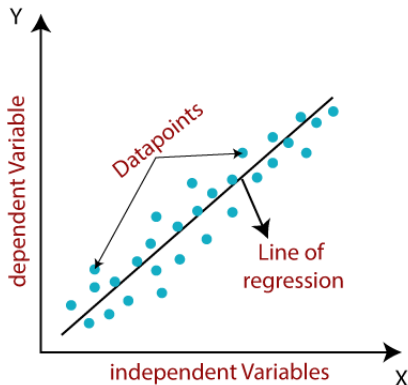
which shows that about 15 decimal digits will be lost. This confirms our observation from the Python code!

- ▶ Hilbert matrix is not the only ill-conditioned matrix, another example is [Vandermonde matrix](#), and many other matrices.
- ▶ If A is singular, even non-square, its condition number can be defined via SVD (later in the course)
- ▶ **Lesson: avoid ill-conditioned matrices in your algorithms!**

Regression analysis and least squares

Regression analysis

- ▶ Build regression model from (large) data sets
- ▶ Finding trend-lines in data
- ▶ Lots of regression in big data and machine learning areas
- ▶ The idea is to minimize distance between data points and the model, for example a polynomial



Example: Linear least squares

- ▶ Ansatz: $y = a_0 + a_1x$
- ▶ y : dependent variable, predicted/measured outcome
- ▶ x : independent variable
- ▶ a_0, a_1 : coefficients to be determined
- ▶ Approximate equations on m points:

$$\begin{array}{rcl} y_1 & \approx & a_0 + a_1x_1 \\ y_2 & \approx & a_0 + a_1x_2 \\ & \vdots & \\ y_m & \approx & a_0 + a_1x_m \end{array} \quad \Rightarrow \quad \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_m \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \end{bmatrix} \approx \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} \quad \Rightarrow \quad Aa \approx y$$

Example: Quadratic least squares

- ▶ Ansatz: $y = a_0 + a_1x + a_2x^2$
- ▶ Approximate equations on m points:

$$\begin{array}{rcl} y_1 & \approx & a_0 + a_1x_1 + a_2x_1^2 \\ y_2 & \approx & a_0 + a_1x_2 + a_2x_2^2 \\ & \vdots & \\ y_m & \approx & a_0 + a_1x_m + a_2x_m^2 \end{array} \quad \Rightarrow \quad \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_m & x_m^2 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} \approx \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

$$Aa \approx y$$

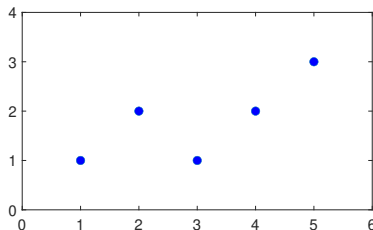
- ▶ The error is $e = Aa - y$
- ▶ in order to find the least squares solution (coefficient vector a) we minimize the 2-norm of the error:

$$\min_{a \in \mathbb{R}^3} \|Aa - y\|_2$$

We will see how!

Example: Quadratic least squares

x	x_1	x_2	x_3	x_4	x_5
y	1	2	1	2	3
	y_1	y_2	y_3	y_4	y_5



$$\begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ 1 & x_3 & x_3^2 \\ 1 & x_4 & x_4^2 \\ 1 & x_5 & x_5^2 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} \approx \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ y_5 \end{bmatrix} \implies \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 4 \\ 1 & 3 & 9 \\ 1 & 4 & 16 \\ 1 & 5 & 25 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} \approx \begin{bmatrix} 1 \\ 2 \\ 1 \\ 2 \\ 3 \end{bmatrix}$$

Overdetermined system of equations (5 equations, 3 unknowns) \implies no exact solution

Least squares fitting (Normal Equations)

The least squares problem leads to **normal equations** (why?):

$$\min_{\mathbf{a} \in \mathbb{R}^3} \|\mathbf{A}\mathbf{a} - \mathbf{y}\|_2 \implies \mathbf{A}^T \mathbf{A} \mathbf{a} = \mathbf{A}^T \mathbf{y}$$

If \mathbf{A} is a full rank matrix then $\mathbf{A}^T \mathbf{A}$ is non-singular and the normal system has a unique solution.

Again the Example:

$$\mathbf{A}^T \mathbf{A} \mathbf{a} = \mathbf{A}^T \mathbf{y} \implies \begin{bmatrix} 5 & 15 & 55 \\ 15 & 55 & 225 \\ 55 & 225 & 979 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} = \begin{bmatrix} 9 \\ 31 \\ 125 \end{bmatrix}$$

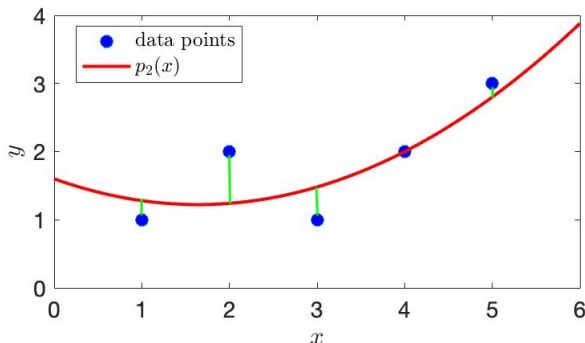
Solve with Gaussian elimination yields (rounded to two decimals)

$$\mathbf{a} = \begin{bmatrix} 1.60 \\ -0.46 \\ 0.14 \end{bmatrix}$$

Least squares fitting (Normal Equations)

Plug solution into the ansatz:

$$\begin{aligned}p_2(x) &= a_0 + a_1x + a_2x^2 \\ &= 1.6 - 0.46x + 0.14x^2\end{aligned}$$



We minimize the sum of squares of vertical distances (green vertical lines), so the method is called **least squares**

Least squares fitting via normal equations

Steps for a polynomial ansatz are:

- ▶ Substitute the data $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$, into the ansatz and form matrix A and vector value \mathbf{y}
- ▶ Solve the normal equations $A^T A \mathbf{a} = A^T \mathbf{y}$ (as A overdetermined)
- ▶ Substitute the solution \mathbf{a} into the ansatz to get the final least squares polynomial

In principle you can use any polynomial (or other function) as ansatz. But some are better than others.

Least squares fitting via normal equations

Problems with normal system:

- **Significant digits may be lost** when A^T is multiplied by A .

Example:

$$A = \begin{bmatrix} 1 & 1 \\ \epsilon & 0 \\ 0 & \epsilon \end{bmatrix} \implies A^T A = \begin{bmatrix} 1 + \epsilon^2 & 1 \\ 1 & 1 + \epsilon^2 \end{bmatrix}$$

In the double precision arithmetic if $\epsilon < 10^{-8}$ then $\text{fl}(1 + \epsilon^2) = 1$, i.e.

$$A^T A = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

which is singular.

- **The condition number** is increased:

$$\text{cond}_2(A^T A) = [\text{cond}_2(A)]^2$$

In our example: $\text{cond}_2(A^T A) = 0.7378 \cdot 10^4$ while $\text{cond}_2(A) = 0.8589 \cdot 10^2$.

Of course 10^4 is a large cond. number for a 3×3 matrix!

Shifting to improve the conditioning

Shift the ansatz by the **mean of data**: $p_2(x) = a_0 + a_1(x - \bar{x}) + a_2(x - \bar{x})^2$

The same example:

x	1	2	3	4	5
y	1	2	1	2	3

Since $\bar{x} = 3$ we get

$$\begin{bmatrix} 1 & x_1 - 3 & (x_1 - 3)^2 \\ 1 & x_2 - 3 & (x_2 - 3)^2 \\ 1 & x_3 - 3 & (x_3 - 3)^2 \\ 1 & x_4 - 3 & (x_4 - 3)^2 \\ 1 & x_5 - 3 & (x_5 - 3)^2 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} \approx \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ y_5 \end{bmatrix} \implies \begin{bmatrix} 1 & -2 & 4 \\ 1 & -1 & 1 \\ 1 & 0 & 0 \\ 1 & 1 & 11 \\ 1 & 2 & 4 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} \approx \begin{bmatrix} 1 \\ 2 \\ 1 \\ 2 \\ 3 \end{bmatrix}$$

Shifting to improve the conditioning

Normal equations become (after shift):

$$A^T A \mathbf{a} = A^T \mathbf{y} \implies \begin{bmatrix} 5 & 0 & 10 \\ 0 & 10 & 0 \\ 10 & 0 & 34 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} = \begin{bmatrix} 9 \\ 4 \\ 20 \end{bmatrix}$$

Solve with Gaussian elimination yields (rounded to two decimals)

$$\mathbf{a} = \begin{bmatrix} 1.51 \\ 0.40 \\ 0.14 \end{bmatrix}$$

Plug solution into the ansatz:

$$\begin{aligned} p_2(x) &= a_0 + a_1(x - 3) + a_2(x - 3)^2 \\ &= 1.51 + 0.4(x - 3) + 0.14(x - 3)^2 \\ &= 1.6 - 0.46x + 0.14x^2 \end{aligned}$$

Solutions are the same but $\text{cond}_2(A^T A) = 19.7$ (compare $\text{cond}_2(A^T A) = 7377.7$ for other ansatz)

Shift and scale to improve the conditioning

Shift and scale the ansatz by the **mean and standard deviation of data**:

$$p_2(x) = a_0 + a_1 \frac{x - \bar{x}}{\sigma} + a_2 \left(\frac{x - \bar{x}}{\sigma} \right)^2$$

The same example:

x	1	2	3	4	5
y	1	2	1	2	3

Since $\bar{x} = 3$ and $\sigma = 1.5811$ we get the normal equations

$$A^T A \mathbf{a} = A^T \mathbf{y} \implies \begin{bmatrix} 5 & 0 & 4 \\ 0 & 4 & 0 \\ 4 & 0 & 5.44 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} = \begin{bmatrix} 9 \\ 2.5298 \\ 12.6491 \end{bmatrix}$$

- ▶ $\text{cond}_2(A^T A) = 7.6$
- ▶ Same polynomial, but even lower cond. number
- ▶ More stable and accurate solution for many problems
- ▶ **There might be still some stability issues for large scale problems even with shifted and scaled ansatz.**

Previous example via normal equation:

```
# Quadratic fitting with shift and scale
import numpy as np

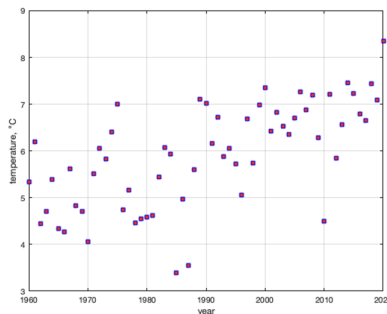
x = np.array([1,2,3,4,5])
y = np.array([1,2,1,2,3])
mu, s = np.mean(x), np.std(x)
x = (x-mu)/s
A = np.ones((len(x),3))
A[:,1] = x; A[:,2] = x**2
a = np.linalg.solve(A.T@A, A.T@y)
print('a = ', a)
print("The model: y = %f+%f(x-%f)/%f+%f((x-%f)/%f)^2"
      % (a[0],a[1],mu,s,a[2],mu,s))
print('cond(ATA) = ', np.linalg.cond(A.T@A))
```

Output:

```
a = [1.51428571 0.56568542 0.28571429]
The model: y = 1.51428+0.565685(x-3)/1.414214+0.285714((x-3)/1.414214)^2
cond(ATA) = 8.293712447937585
```

Shift and scale to improve the conditioning

Data from Lab exercise 2 (temperature data):



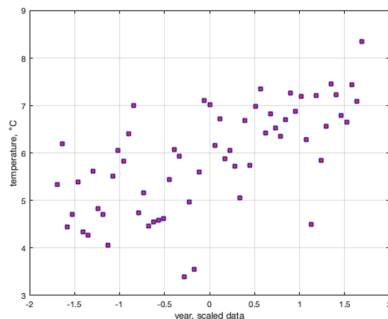
Original data

Polynomial degree 1:

$$\text{cond}_2(A^T A) = 5 \times 10^{10}$$

Polynomial degree 2:

$$\text{cond}_2(A^T A) = 5 \times 10^{21}$$



Centered and scaled data

Polynomial degree 1:

$$\text{cond}_2(A^T A) = 1.02$$

Polynomial degree 2:

$$\text{cond}_2(A^T A) = 7.67$$

Least squares fitting

Data from Lab exercise 2 (temperature data): What happens in Python when running the case Polynomial degree 2 when $\text{cond}_2(A^T A) = 5 \times 10^{21}$?

- ▶ If `numpy.linalg.solve` is used as equation solver, nothing happens.
- ▶ If `scipy.linalg.solve` is used, get a warning:

```
v = scipy.linalg.solve(ATA, ATy)
```

```
LinAlgWarning: Ill-conditioned matrix (rcond=3.1176e-22):  
result may not be accurate.
```

- ▶ The difference between the `Polynomial.fit` solution and normal equation solution is approximately 10^{-7} (the norm of the difference). Interpretation: Loose accuracy, roughly 9 decimal places (remember the best we can expect is 10^{-16})