Models and Languages for Computational Systems Biology: Stochastic Simulation

Stephen Gilmore

Laboratory for Foundations of Computer Science
School of Informatics
University of Edinburgh

February 12, 2007

Stochastic: Propensity function

As explicitly derived by Gillespie, the stochastic model uses basic Newtonian physics and thermodynamics to arrive at a form often termed the propensity function that gives the probability a_μ of reaction μ occurring in time interval $(t, t + \mathrm{d}t)$.

$$a_\mu \mathrm{d}t = h_\mu c_\mu \mathrm{d}t$$

where the M reaction mechanisms are given an arbitrary index μ $(1 \leq \mu \leq M)$, h_μ denotes the number of possible combinations of reactant molecules involved in reaction μ, and c_μ is a stochastic rate constant.

Stochastic: Fundamental hypothesis

The rate constant c_μ is dependent on the radii of the molecules involved in the reaction, and their average relative velocities – a property that is itself a direct function of the temperature of the system and the individual molecular masses.

These quantities are basic chemical properties which for most systems are either well known or easily measurable. Thus, for a given chemical system, the propensity functions, a_μ can be easily determined.

Stochastic: Grand probability function

The stochastic formulation proceeds by considering the grand probability function $\Pr(\mathbf{X}; t) \equiv$ probability that there will be present in the volume V at time t, X_i of species S_i, where $\mathbf{X} \equiv (X_1, X_2, \ldots X_N)$ is a vector of molecular species populations.

Evidently, knowledge of this function provides a complete understanding of the probability distribution of all possible states at all times.

Stochastic: Infinitesimal time interval

By considering a discrete infinitesimal time interval $(t, t + \mathrm{d}t)$ in which either 0 or 1 reactions occur we see that there exist only $M + 1$ distinct configurations at time t that can lead to the state $\mathbf{X}$ at time $t + \mathrm{d}t$.

$$\begin{aligned}\Pr(\mathbf{X}; t + \mathrm{d}t) &= \Pr(\mathbf{X}; t) \Pr(\text{no state change over } \mathrm{d}t) \\ &+ \textstyle\sum_{\mu=1}^{M} \Pr(\mathbf{X} - \mathbf{v}_\mu; t) \Pr(\text{state change to } \mathbf{X} \text{ over } \mathrm{d}t)\end{aligned}$$

where $\mathbf{v}_\mu$ is a stoichiometric vector defining the result of reaction μ on state vector $\mathbf{X}$, i.e. $\mathbf{X} \to \mathbf{X} + \mathbf{v}_\mu$ after an occurrence of reaction μ.

Stochastic: State change probabilities

$\Pr(\text{no state change over } \mathrm{d}t)$

$$1 - \sum_{\mu=1}^{M} a_\mu(\mathbf{X}) \mathrm{d}t$$

$\Pr(\text{state change to } \mathbf{X} \text{ over } \mathrm{d}t)$

$$\sum_{\mu=1}^{M} \Pr(\mathbf{X} - \mathbf{v}_\mu; t) a_\mu(\mathbf{X} - \mathbf{v}_\mu) \mathrm{d}t$$

Stochastic: Partial derivatives

$$\frac{\partial \Pr(\mathbf{X}; t)}{\partial t} = \lim_{\mathrm{d}t \to 0} \frac{\Pr(\mathbf{X}; t + \mathrm{d}t) - \Pr(\mathbf{X}; t)}{\mathrm{d}t}$$

Stochastic: Chemical Master Equation

Applying this, and re-arranging the former, leads us to an important *partial differential equation* (PDE) known as the Chemical Master Equation (CME).

$$\frac{\partial \Pr(\mathbf{X}; t)}{\partial t} = \sum_{\mu=1}^{M} a_\mu(\mathbf{X} - \mathbf{v}_\mu) \Pr(\mathbf{X} - \mathbf{v}_\mu; t) - a_\mu(\mathbf{X}) \Pr(\mathbf{X}; t)$$

The problem with the Chemical Master Equation

- The CME is really a set of nearly as many coupled ordinary differential equations as there are combinations of molecules that can exist in the system!
- The CME can be solved analytically for only a very few very simple systems, and numerical solutions are usually prohibitively difficult.

D. Gillespie and L. Petzold.
chapter *Numerical Simulation for Biochemical Kinetics*, in *System Modelling in Cellular Biology*, editors Z. Szallasi, J. Stelling and V. Periwal.
MIT Press, 2006.

Stochastic simulation algorithms

Gillespie's Stochastic Simulation Algorithm (SSA) is essentially an exact procedure for numerically simulating the time evolution of a well-stirred chemically reacting system by taking proper account of the randomness inherent in such a system.

It is rigorously based on the same microphysical premise that underlies the chemical master equation and gives a more realistic representation of a system's evolution than the deterministic reaction rate equation (RRE) represented mathematically by ODEs.

As with the chemical master equation, the SSA converges, in the limit of large numbers of reactants, to the same solution as the law of mass action.

Gillespie's exact SSA (1977)

- The algorithm takes time steps of variable length, based on the rate constants and population size of each chemical species.
- The probability of one reaction occurring relative to another is dictated by their relative propensity functions.
- According to the correct probability distribution derived from the statistical thermodynamics theory, a random variable is then used to choose which reaction will occur, and another random variable determines how long the step will last.
- The chemical populations are altered according to the stoichiometry of the reaction and the process is repeated.

Stochastic simulation: realisations and ensembles

The SSA computes one realisation of a dynamic trajectory of a chemically reacting system. Often an ensemble of trajectories is computed, to obtain an estimate of the probability density function of the system.

The dynamic evolution of the probability density function is given by the Chemical Master Equation.

Gillespie's SSA is a Monte Carlo Markov Chain simulation

The SSA is a Monte Carlo type method. With the SSA one may approximate any variable of interest by generating many trajectories and observing the statistics of the values of the variable. Since many trajectories are needed to obtain a reasonable approximation, the efficiency of the SSA is of critical importance.

Computational cost of Gillespie's exact algorithm

The cost of this detailed stochastic simulation algorithm is the likely large amounts of computing time.

The key issue is that the time step for the next reaction can be very small indeed if we are to guarantee that only one reaction can take place in a given time interval.

Increasing the molecular population or number of reaction mechanisms necessarily requires a corresponding decrease in the time interval. The SSA can be very computationally inefficient especially when there are large numbers of molecules or the propensity functions are large.

Gibson and Bruck (2000)

Gibson and Bruck refined the first reaction SSA of Gillespie by reducing the number of random variables that need to be simulated.

This can be effective for systems in which some reactions occur much more frequently than others.

M.A. Gibson and J. Bruck.
Efficient exact stochastic simulation of chemical systems with many species and many channels.
J. Comp. Phys., 104:1876–1889, 2000.

Variants of SSA

Gillespie developed two different but equivalent formulations of the SSA: the Direct Method (DM) and the First Reaction Method (FRM). A third formulation of the SSA is the Next Reaction Method (NRM) of Gibson and Bruck. The NRM can be viewed as an extension of the FRM, but it is much more efficient than the latter.

It was widely believed that Gibson and Bruck's method (the Next Reaction Method) was more efficient than Gillespie's Direct Method (DM). This conclusion is based on a count of arithmetic operations.

Gibson and Bruck challenged (2004)

It was established by Cao, Li and Petzold (2004) that Gibson and Bruck's analysis misses the dominant cost of the NRM, which is maintaining the priority queue data structure of the tentative reaction times and that good implementations of DM such as the Optimised Direct Method (ODM) have lower asymptotic complexity than Gibson and Bruck's method.

Y. Cao, H. Li, and L. Petzold.
Efficient formulation of the stochastic simulation algorithm for chemically reacting systems.
J. Chem. Phys, 121(9):4059–4067, 2004.

Enhanced stochastic simulation techniques

If the system under study possesses a macroscopically infinitesimal timescale so that during any dt all of the reaction channels can fire many times, yet none of the propensity functions change appreciably, then the discrete Markov process as described by the SSA can be approximated by a continuous Markov process.

This Markov process is described by the Chemical Langevin Equation (CLE), which is a stochastic ordinary differential equation (SDE).

Stochastic Differential Equations

A stochastic differential equation (SDE)

$$dX_t = a(t, X_t)dt + b(t, X_t)dW_t$$

is interpreted as a stochastic integral equation

$$X_t = X_{t_0} + \int_{t_0}^{t} a(s, X_s)ds + \int_{t_0}^{t} b(s, X_s)dW_s$$

where the first integral is a Lebesgue (or Riemann) integral for each sample path and the second integral is usually an Ito integral.

Chemical Langevin Equation

The Langevin equation

$$dX_t = -aX_t dt + dW_t$$

is a linear SDE with additive noise. The solution for $t_0 = 0$ is

$$X_t = X_0 e^{-at} + e^{-at} \int_0^t e^{as} dW_s$$

Gillespie's tau-leap method (2001)

Gillespie proposed two new methods, namely the τ-leap method and the midpoint τ-leap method in order to improve the efficiency of the SSA while maintaining acceptable losses in accuracy.

Daniel T. Gillespie.
Approximate accelerated stochastic simulation of chemically reacting systems.
J. Comp. Phys., 115(4):1716–1733, 2001.

The key idea here is to take a larger time step and allow for more reactions to take place in that step, but under the proviso that the propensity functions do not change too much in that interval. By means of a Poisson approximation, the tau-leaping method can "leap over" many reactions.

Gillespie's tau-leap method

For many problems, the tau-leaping method can approximate the stochastic behaviour of the system very well.

The tau-leaping method connects the SSA in the discrete stochastic regime to the explicit Euler method for the chemical Langevin equation in the continuous stochastic regime and the RRE in the continuous deterministic regime.

However, the use of approximation in Poisson methods leads to the possibility of negative molecular numbers being predicted — something with no physical explanation.

Gillespie's Modified Poisson tau-leap methods (2005)

Gillespie's modified Poisson tau-leaping method introduces a second control parameter whose value dials the procedure from the original Poisson tau-leaping method at one extreme to the exact SSA at the other.

Any reaction channel with a positive propensity function which is within n_c firings of exhausting its reactants is termed a *critical* reaction.

Y. Cao, D. Gillespie, and L. Petzold.
Avoiding negative populations in explicit tau leaping.
J. Chem. Phys, 123(054104), 2005.

Gillespie's Modified Poisson tau-leap methods (2006)

The modified algorithm chooses τ in such a way that no more than *one* firing of *all* the critical reactions can occur during the leap. The probability of producing a negative population is reduced to nearly zero.

If a negative population *does* occur the leap can simply be rejected and repeated with τ reduced by half, or the entire simulation can be abandoned and repeated for larger n_c.

Y. Cao, D. Gillespie, and L. Petzold.
Efficient stepsize selection for the tau-leaping method.
J. Chem. Phys, 2006.
To appear.

Family of stochastic simulation algorithms

FASTEST, BEST	
Discrete, exact	**Continuous, approximate**
	Modified Poisson τ leap (2005)
	τ leap (2001)
Logarithmic Direct Method (2006)	
Sorting Direct Method (2005)	
Optimised Direct Method (2004)	
Next Reaction Method (2000)	
Direct Method (1977)	
First Reaction Method (1977)	
SLOWEST, WORST	