



CHAPTER 7

More SQL: Complex Queries, Triggers, Views, and Schema Modification

Chapter 7 Outline

- More Complex SQL Retrieval Queries
- Specifying Semantic Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Modification in SQL

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping

Comparisons Involving NULL and Three-Valued Logic

- Meanings of NULL
 - **Unknown value**
 - **Unavailable or withheld value**
 - **Not applicable attribute**
- Each individual NULL value considered to be different from every other NULL value
- SQL uses a three-valued logic:
 - TRUE, FALSE, and UNKNOWN (like Maybe)
- **NULL = NULL comparison is avoided**

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- SQL allows queries that check whether an attribute value is NULL
 - IS or IS NOT NULL

Query 18. Retrieve the names of all employees who do not have supervisors.

```
Q18:  SELECT  Fname, Lname
      FROM    EMPLOYEE
      WHERE   Super_ssn IS NULL;
```

Nested Queries, Tuples, and Set/Multiset Comparisons

- **Nested queries**

- Complete select-from-where blocks within WHERE clause of another query
- **Outer query and nested subqueries**

- **Comparison operator** `IN`

- Compares value v with a set (or multiset) of values V
- Evaluates to `TRUE` if v is one of the elements in V

Nested Queries (cont'd.)

Q4A: **SELECT** **DISTINCT** Pnumber
 FROM **PROJECT**
 WHERE Pnumber **IN**
 (**SELECT** Pnumber
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND**
 Mgr_ssn=Ssn **AND** Lname='Smith')

 OR
 Pnumber **IN**
 (**SELECT** Pno
 FROM WORKS_ON, EMPLOYEE
 WHERE Essn=Ssn **AND** Lname='Smith');

{ }

{1, 2}

{1, 2}

Nested Queries (cont'd.)

- Avoid potential errors and ambiguities
 - Create tuple variables (aliases) for all tables referenced in SQL query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT    E.Fname, E.Lname
      FROM      EMPLOYEE AS E
      WHERE     E.Ssn IN ( SELECT    Essn
                          FROM      DEPENDENT AS D
                          WHERE     E.Fname=D.Dependent_name
                          AND E.Sex=D.Sex );
```

Specifying Joined Tables in the FROM Clause of SQL

■ Joined table

- Permits users to specify a table resulting from a join operation in the FROM clause of a query

■ The FROM clause in Q1A

- Contains a single joined table. JOIN may also be called INNER JOIN

Q1A: SELECT Fname, Lname, Address
 FROM (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
 WHERE Dname='Research';

FName	MInt	LName	Dno	Address	Dname	Dnumber
John	B	Smith	5	731 Fondern,..		Research	5	
Fraklin	T	Wong	5	638 Voss,..		Research	5	
Alicia	J	Zelaya	4	3321 Castle,..		Administration	4	
..			..					

Join

Different Types of JOINed Tables in SQL

- Specify different types of join
 - NATURAL JOIN
 - Various types of OUTER JOIN (LEFT, RIGHT, FULL)
- NATURAL JOIN on two relations R and S
 - No join condition specified
 - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

```
Q1B:      SELECT      Fname, Lname, Address
           FROM        (EMPLOYEE NATURAL JOIN
                        (DEPARTMENT AS DEPT (Dname, Dno, Mssn,
                                             Msdate)))
           WHERE       Dname='Research';
```

The above works with $EMPLOYEE.Dno = DEPT.Dno$ as an implicit join condition

INNER and OUTER Joins

■ INNER JOIN

- Default type of join in a joined table
- Tuple is included in the result only if a matching tuple exists in the other relation

Students

Student_ID	Student_Name	Advisor_ID
1	Student_1	1
2	Student_2	8
4	Student_4	2
5	Student_5	3
7	Student_7	3
9	Student_9	1
10	Student_10	3

Advisors


Advisor_ID	Advisor_Name
1	Advisor 1
3	Advisor 3
5	Advisor 5

```
select Student_Name,  
       Advisor_Name  
from Students S INNER JOIN  
     Advisors A ON  
     S.Advisor_ID=A.Advisor_ID
```

Block Name Join Block

Join Type

[Show Output](#) Block output (Inner Join) 5 rows.



Student_Name	Advisor_Name
Student_1	Advisor 1
Student_5	Advisor 3
Student_7	Advisor 3
Student_9	Advisor 1
Student_10	Advisor 3

INNER and OUTER Joins

■ LEFT OUTER JOIN

- Every tuple in left table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of right table

Students

Advisors

Student_ID	Student_Name	Advisor_ID
1	Student_1	1
2	Student_2	8
4	Student_4	2
5	Student_5	3
7	Student_7	3
9	Student_9	1
10	Student_10	3


Advisor_ID	Advisor_Name
1	Advisor 1
3	Advisor 3
5	Advisor 5

```
select *  
from  
Students S  
LEFT OUTER JOIN Advisors A  
ON S.Advisor_ID=A.Advisor_ID
```

Block Name: Join Block 2

Join Type: Outer Join on Input 1

Block output (Outer Join on Input 1) 7 rows.



Student_Name	Advisor_Name
Student_2	_null_
Student_4	_null_
Student_1	Advisor 1
Student_5	Advisor 3
Student_7	Advisor 3
Student_9	Advisor 1
Student_10	Advisor 3

INNER and OUTER Joins

- LEFT OUTER JOIN
 - Every tuple in left table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of right table
- RIGHT OUTER JOIN
 - Every tuple in right table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of left table

Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
 - **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- **Grouping**
 - Create subgroups of tuples before summarizing
- To select entire groups, **HAVING** clause is used
- Aggregate functions can be used in the **SELECT** clause or in a **HAVING** clause

Renaming Results of Aggregation

- Following query returns a single row of computed values from EMPLOYEE table:

Q19:

```
SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG  
       (Salary)  
FROM   EMPLOYEE;
```

- The result can be presented with new names:

```
Q19A:      SELECT    SUM (Salary) AS Total_Sal, MAX (Salary) AS  
                Highest_Sal, MIN (Salary) AS Lowest_Sal, AVG  
                (Salary) AS Average_Sal  
  
            FROM      EMPLOYEE;
```

Aggregate functions

- For example: get all employees whose salary is $>$ than 30.
- Some advanced operations may address sets of tuples.
- For example: how many employees have a salary $>$ than 30?
- SQL provides this functionality through aggregate functions.

```
select *  
from Employee  
where Salary > 30
```

<i>Name</i>	<i>Surname</i>	<i>Department</i>	<i>Supervisor</i>	<i>Salary</i>
John	White	1	2	36
Mark	Frank	1	3	46
Moan	Jones	2	1	27

Example

- Select the number of employees working at Department number 1.

Evaluating aggregate queries (1)

```
select      *  
from Employee  
where Department = 1
```

<i>Name</i>	<i>Surname</i>	<i>Department</i>	<i>Supervisor</i>	<i>Salary</i>
John	White	1	2	36
Mark	Frank	1	3	46

Evaluating aggregate queries (2)

```
select count(*) AS numberOfEmployees
from Employee
where Department = 1
```

	<i>Name</i>	<i>Surname</i>	<i>Department</i>	<i>Supervisor</i>	<i>Salary</i>
→	John	White	1	2	36
→	Mark	Frank	1	3	46


<i>numberOfEmployees</i>
2

Other standard aggregate functions

- count, sum, max, min, avg.
- Check the manual of the system you want to use for other options.

Target of the aggregate function

select count(distinct salary) AS numOfDistinctSalaries
from Employee



Aggregate Functions in SQL (cont'd.)

- NULL values are discarded when aggregate functions are applied to a particular column

Query 20. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:  SELECT    SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
      FROM      (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE     Dname='Research';
```

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:  SELECT    COUNT (*)
      FROM      EMPLOYEE;
```

```
Q22:  SELECT    COUNT (*)
      FROM      EMPLOYEE, DEPARTMENT
      WHERE     DNO=DNUMBER AND DNAME='Research';
```


Grouping: The GROUP BY Clause

- **Partition** relation into subsets of tuples
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- **GROUP BY** clause
 - Specifies grouping attributes
- **COUNT (*)** counts the number of rows in the group

Sum of salaries per department

select Department, Salary
from Employee

<i>Department</i>	<i>Salary</i>
1	45
2	36
1	40
3	45
4	80
4	73
1	40
2	46

Sum of salaries per department

```
select Department,      Salary
from Employee
GROUP BY Department
```

<i>Department</i>	<i>Salary</i>
1	45
1	40
1	40
2	36
2	46
3	45
4	80
4	73

Sum of salaries per department

```
select Department, sum(Salary)
from Employee
GROUP BY Department
```

<i>Department</i>	
1	125
2	82
3	45
4	153

<i>Department</i>	<i>Salary</i>
1	45
1	40
1	40
2	36
2	46
3	45
4	80
4	73

Sum of salaries per department

```
select Department, sum(Salary) AS allSalary  
from Employee  
GROUP BY Department
```

<i>Department</i>	<i>allSalary</i>
1	125
2	82
3	45
4	153

Predicates on groups

- **HAVING** clause

- Provides a condition to select or reject an entire group:

```
select Department, sum(Salary)
from Employee
group by Department
HAVING sum(Salary) > 100
```

EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.
- **INCORRECT QUERY:**

```
SELECT      Dno, COUNT (*)  
FROM        EMPLOYEE  
WHERE       Salary>40000  
GROUP BY    Dno  
HAVING      COUNT (*) > 5;
```


Combining the WHERE and the HAVING Clause (continued)

Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

Query 28. For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28:  SELECT  Dnumber, COUNT (*)
      FROM    DEPARTMENT, EMPLOYEE
      WHERE   Dnumber=Dno AND Salary>40000 AND
            ( SELECT  Dno
              FROM    EMPLOYEE
              GROUP BY Dno
              HAVING  COUNT (*) > 5)
```

Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table derived from other tables called the **defining tables**
 - Considered to be a virtual table that is not necessarily populated

Specification of Views in SQL

■ **CREATE VIEW** command

- **View** – a single table derived from other tables.
- The tables can be **base tables** or other previously defined views.
 - Give table name, list of attribute names, and a query to specify the contents of the view
 - In V1, attributes retain the names from base tables. In V2, attributes are assigned names

```
V1:  CREATE VIEW  WORKS_ON1
      AS SELECT   Fname, Lname, Pname, Hours
      FROM        EMPLOYEE, PROJECT, WORKS_ON
      WHERE       Ssn=Essn AND Pno=Pnumber;
```

```
V2:  CREATE VIEW  DEPT_INFO(Dept_name, No_of_emps, Total_sal)
      AS SELECT   Dname, COUNT (*), SUM (Salary)
      FROM        DEPARTMENT, EMPLOYEE
      WHERE       Dnumber=Dno
      GROUP BY    Dname;
```

Specification of Views in SQL (cont'd.)

- Once a View is defined, SQL queries can use the View relation in the FROM clause
- View is always up-to-date
 - Responsibility of the DBMS and not the user
- **DROP VIEW** command
 - Dispose of a view

The DROP Command

- DROP command
 - Used to drop named schema elements, such as tables, domains, or constraint
- Drop behavior options:
 - CASCADE and RESTRICT
- Example:
 - DROP SCHEMA COMPANY CASCADE;
 - This removes the schema and all its elements including tables, views, constraints, etc.

The ALTER table command

- **Alter table actions** include:
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD
COLUMN Job VARCHAR(12) ;`

Adding and Dropping Constraints

- Change constraints specified on a table
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```

Summary

- Complex SQL:
 - Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- Handling semantic constraints with CREATE ASSERTION and CREATE TRIGGER
- CREATE VIEW statement and materialization strategies
- Schema Modification for the DBAs using ALTER TABLE , ADD and DROP COLUMN, ALTER CONSTRAINT etc.